

6.S081 2020 Lecture 1: O/S overview

Overview

- * 6.S081 goals
 - * Understand operating system (O/S) design and implementation
 - * Hands-on experience extending a small O/S
 - * Hands-on experience writing systems software
- * What is the purpose of an O/S?
 - * Abstract the hardware for convenience and portability
 - * Multiplex the hardware among many applications
 - * Isolate applications in order to contain bugs
 - * Allow sharing among cooperating applications
 - * Control sharing for security
 - * Don't get in the way of high performance
 - * Support a wide range of applications
- * Organization: layered picture
[user/kernel diagram]
 - user applications: vi, gcc, DB, &c
 - kernel services
 - h/w: CPU, RAM, disk, net, &c
 - * we care a lot about the interfaces and internal kernel structure

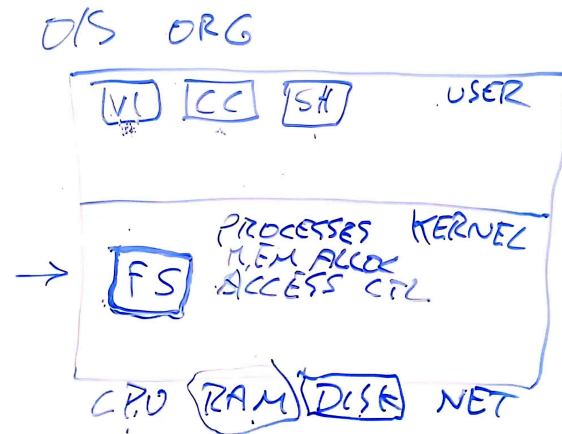
- * What services does an O/S kernel typically provide?
 - * process (a running program)
 - * memory allocation
 - * file contents
 - * file names, directories
 - * access control (security)
 - * many others: users, IPC, network, time, terminals
- * What's the application / kernel interface?
 - * "System calls"
 - * Examples, in C, from UNIX (e.g. Linux, macOS, FreeBSD):

```
fd = open("out", 1);
write(fd, "hello\n", 6);
pid = fork();
```

- * These look like function calls but they aren't
- * Why is O/S design+implementation hard and interesting?
 - * unforgiving environment: quirky h/w, hard to debug
 - * many design tensions:
 - efficient vs abstract/portable/general-purpose
 - powerful vs simple interfaces
 - flexible vs secure
 - * features interact: `fd = open(); fork()``
 - * uses are varied: laptops, smart-phones, cloud, virtual machines, embedded
 - * evolving hardware: NVRAM, multi-core, fast networks
- * You'll be glad you took this course if you...
 - * care about what goes on under the hood
 - * like infrastructure
 - * need to track down bugs or security problems
 - * care about high performance

Class structure

- * Online course information:
 - <https://pdos.csail.mit.edu/6.828/2020/schedule.html>
 - <https://pdos.csail.mit.edu/6.S081/> -- schedule, assignments, labs
 - Piazza -- announcements, discussion, lab help



- * Lectures
 - * O/S ideas
 - * case study of xv6, a small O/S, via code and xv6 book
 - * lab background
 - * O/S papers
 - * submit a question about each reading, before lecture.
- * Labs:

The point: hands-on experience
Mostly one week each.
Three kinds:

 - Systems programming (due next week...)
 - O/S primitives, e.g. thread switching.
 - O/S kernel extensions to xv6, e.g. network.

Use piazza to ask/answer lab questions.
Discussion is great, but please do not look at others' solutions!
- * Grading:

70% labs, based on tests (the same tests you run).
20% lab check-off meetings: we'll ask you about randomly-selected labs.
10% home-work and class/piazza discussion.
No exams, no quizzes.
Note that most of the grade is from labs. Start them early!

Introduction to UNIX system calls

- * Applications see the O/S via system calls; that interface will be a big focus.
let's start by looking at how programs use system calls.
you'll use these system calls in the first lab.
and extend and improve them in subsequent labs.
- * I'll show some examples, and run them on xv6.
xv6 has similar structure to **UNIX** systems such as Linux.
but much simpler -- you'll be able to digest all of xv6
accompanying book explains how xv6 works, and why
why UNIX?
open source, well documented, clean design, widely used
studying xv6 will help if you ever need to look inside Linux
xv6 has two roles in 6.S081:
example of core functions: virtual memory, multi-core, interrupts, &c
starting point for most of the labs
xv6 runs on RISC-V, as in current 6.004 // copy.c: copy input to output.
you'll run xv6 under the qemu machine emulator
- * example: copy.c, copy input to output
read bytes from input, write them to the output
\$ copy
copy.c is written in C
Kernighan and Ritchie (K&R) book is good for learning C
you can find these example programs via the schedule on the web site
read() and write() are system calls
first read()/write() **argument** is a "file descriptor" (fd)
passed to kernel to tell it which "open file" to read/write
must previously have been opened
an FD connects to a file/device/socket/&c
a process can open many files, have many FDs
UNIX convention: fd 0 is "standard input", 1 is "standard output"
second read() argument is a **pointer** to some memory into which to read
third argument is the maximum number of bytes to read exit(0);
read() may read less, but not more
return value: number of bytes actually read, or -1 for error
note: copy.c does not care about the format of the data
UNIX I/O is 8-bit bytes
interpretation is application-specific, e.g. database records, C source, &c
where do file descriptors come from?

```
#include "kernel/types.h"
#include "user/user.h"

int main()
{
    char buf[64];
    while(1){
        int n = read(0, buf, sizeof(buf));
        if(n <= 0)
            break;
        write(1, buf, n);
    }
}
```

// open.c: create a file, write to it.

* example: open.c, create a file

```
$ open
```

```
$ cat output.txt
```

open() creates a file, returns a file descriptor (or -1 for error)

FD is a small integer

FD indexes into a per-process table maintained by kernel

[user/kernel diagram]

different processes have different FD name-spaces

i.e. FD 1 often means different things to different processes

these examples ignore errors -- don't be this sloppy!

Figure 1.2 in the xv6 book lists system call arguments/return

or look at UNIX man pages, e.g. "man 2 open"

#include "kernel/types.h"

#include "user/user.h"

#include "kernel/fcntl.h"

int main()

```
{
    int fd = open("output.txt", O_WRONLY |
O_CREATE);
```

```
    write(fd, "ooo\n", 4);
```

```
    exit(0);
}
```

* what happens when a program calls a system call like open()?

looks like a function call, but it's actually a special instruction

hardware saves some user registers

hardware increases privilege level

hardware jumps to a known "entry point" in the kernel

now running C code in the kernel

kernel calls system call implementation

open() looks up name in file system

it might wait for the disk

it updates kernel data structures (cache, FD table)

restore user registers

reduce privilege level

jump back to calling point in the program, which resumes

we'll see more detail later in the course

* I've been typing to UNIX's command-line interface, the shell.

the shell prints the "\$" prompts.

the shell lets you run UNIX command-line utilities

useful for system management, messing with files, development, scripting

```
$ ls // fork.c: create a new process
```

```
$ ls > out
```

```
$ grep x < out cat out
```

UNIX supports other styles of interaction too

window systems, GUIs, servers, routers, &c.

but time-sharing via the shell was the original focus of UNIX.

we can exercise many system calls via the shell.

#include "kernel/types.h"

#include "user/user.h"

int main()

```
{
    int pid;
    pid = fork();
    printf("fork() returned %d\n", pid);
```

* example: fork.c, create a new process

the shell creates a new process for each command you type, e.g. for

```
$ echo hello
```

the fork() system call creates a new process

```
$ fork
```

the kernel makes a copy of the calling process

instructions, data, registers, file descriptors, current directory

"parent" and "child" processes

only difference: fork() returns a pid in parent, 0 in child

a pid (process ID) is an integer, kernel gives each process a different pid

thus:

fork.c's "fork() returned" executes in *both* processes

the "if(pid == 0)" allows code to distinguish

ok, fork lets us create a new process

how can we run a program in that process?

if(pid == 0){

```
    printf("child\n");
```

```
} else {
```

```
    printf("parent\n");
}
```

```
exit(0);
```

// exec.c: replace a process with an executable file

#include "kernel/types.h"

#include "user/user.h"

int

main()

```
{
    char*argv[] = { "echo", "this", "is", "echo", 0 };
```

```
    exec("echo", argv);
```

```
    printf("exec failed!\n");
```

```
    exit(0);
}
```

* example: exec.c, replace calling process with an executable file

how does the shell run a program, e.g.

```
$ echo a b c
```

a program is stored in a file: instructions and initial memory

created by the compiler and linker

so there's a file called echo, containing instructions

```
$ exec
```

exec() replaces current process with an executable file

discards instruction and data memory

loads instructions and memory from the file

preserves file descriptors
 exec(filename, argument-array)
 argument-array holds command-line arguments; exec passes to main()
 cat user/echo.c
 echo.c shows how a program looks at its command-line arguments

* example: forkexec.c, fork() a new process, exec() a program

\$ forkexec

forkexec.c contains a common UNIX idiom:

fork() a child process
 exec() a command in the child
 parent wait()s for child to finish
 the shell does fork/exec/wait for every command you type
 after wait(), the shell prints the next prompt
 to run in the background -- & -- the shell skips the wait()
 exit(status) -> wait(&status)

status convention: 0 = success, 1 = command encountered an **error**

note: the fork() copies, but exec() discards the copied memory

this may seem wasteful

you'll transparently eliminate the copy in the "copy-on-write" lab

* example: redirect.c, redirect the output of a command

what does the shell do for this?

\$ echo hello > out

answer: fork, change FD 1 in child, exec echo

\$ redirect

\$ cat output.txt

note: open() always chooses lowest unused FD; 1 due to close(1).

fork, FDs, and exec interact nicely to implement I/O redirection

separate fork-then-exec give child a chance to change FDs before exec

FDs provide indirection

commands just use FDs 0 and 1, don't have to know where they go

exec preserves the FDs that sh set up

thus: only sh has to know about I/O redirection, not each program

* It's worth asking "why" about design decisions:

Why these I/O and process abstractions? Why not something else?

Why provide a file system? Why not let programs use the disk their own way?

Why FDs? Why not pass a filename to write()?

Why are files streams of bytes, not disk blocks or formatted records?

Why not combine fork() and exec()?

The UNIX design works well, but we will see other designs!

* example: pipe1.c, communicate through a pipe

how does the shell implement

\$ ls | grep x

\$ pipe1

an FD can refer to a "pipe", as well as a file

the pipe() system call creates two FDs

read from the first FD

write to the second FD

the kernel maintains a buffer for each pipe

[u/k diagram]

write() appends to the buffer

read() waits until there is data

* example: pipe2.c, communicate between processes

pipes combine well with fork() to implement ls | grep x}

shell creates a pipe,

then forks (twice),

then connects ls's FD 1 to pipe's write FD,

and grep's FD 0 to the pipe

[diagram]

\$ pipe2 -- a simplified version

pipes are a separate abstraction, but combine well w/ fork()

```
int
main()
{
    int pid, status;

    pid = fork();
    if(pid == 0){
        char *argv[] = { "echo", "THIS", "IS",
"ECHO", 0 };
        exec("echo", argv);
        printf("exec failed!\n");
        exit(1);
    } else {
        printf("parent waiting\n");
        wait(&status);
        printf("the child exited with status
%d\n", status);
    }
}

int main()
{
    int pid;

    pid = fork();
    if(pid == 0){
        close(1);
        open("output.txt",
O_WRONLY|O_CREATE);

        char *argv[] = { "echo", "this",
"is", "redirected", "echo", 0 };
        exec("echo", argv);
        printf("exec failed!\n");
        exit(1);
    } else {
        wait((int *) 0);
    }

    exit(0);
}
```

```
int main()
{
    int fds[2];
    char buf[100];
    int n;

    // create a pipe, with two FDs in fds[0], fds[1].
    pipe(fds);

    write(fds[1], "this is pipe1\n", 14);
    n = read(fds[0], buf, sizeof(buf));

    write(1, buf, n);
}
```

```
exit(0);
```

* example: list.c, list files in a directory
 how does ls get a list of the files in a directory?
 you can open a directory and read it -> file names
 "." is a pseudo-name for a process's **current directory**
 see ls.c for more details

* Summary

- * We've looked at UNIX's I/O, file system, and process abstractions.
- * The interfaces are simple -- just integers and I/O buffers.
- * The abstractions combine well, e.g. for I/O redirection.

You'll use these system calls in the first lab, due next week.

```

struct dirent {
    ushort inum;
    char name[14];
};

int
main()
{
    int fd;
    struct dirent e;

    fd = open(".", 0);
    while(read(fd, &e, sizeof(e)) == sizeof(e)){
        if(e.name[0] != '\0'){
            printf("%s\n", e.name);
        }
    }
    exit(0);
}

```

```

int
main()
{
    int n, pid;
    int fds[2];
    char buf[100];

    // create a pipe, with two FDs in fds[0], fds[1].
    pipe(fds);

    pid = fork();
    if (pid == 0) {
        write(fds[1], "this is pipe2\n", 14);
    } else {
        n = read(fds[0], buf, sizeof(buf));
        write(1, buf, n);
    }

    exit(0);
}

```